

Introduction

Business objective: Prediction the buying price of a car given :

- Maintenance
- Number of doors
- Lug Boot Size
- Safety
- Class Value

Assumptions:

- Higher comfort - luggage boot size, number of doors & people, safety; higher pricing (positive correlation)
- Higher class ; higher pricing (positive correlation)

Table content

1. [Loading the data](#)
2. [Discoverng and visualizing the data to gain insights](#)
3. [Preparing the data for machine learning algorithms](#)
4. [Training, evaluating and error analysis on the training set](#)
 - [Logistic Regression](#)
 - [Logistic Regression - Grid Search](#)
 - [K-Nearest Neighbours](#)
 - [Decision Trees](#)
 - [Decision Trees - Grid Search](#)
 - [Multi Layered Perceptron](#)
 - [Multi Layered Perceptron - Grid Search](#)
 - [Ensembling - Random Forest](#)
5. [Selecting the best model](#)
6. [Conclusions: Presenting solution](#)

```
In [1]: import pandas as pd
        from pandas.plotting import scatter_matrix
        import matplotlib.pyplot as plt
        import seaborn as sns
        import numpy as np
        %matplotlib inline
```

```
In [2]: import warnings
        warnings.filterwarnings('ignore')
```

1. Loading the data

```
In [3]:
```

```
ATTRIBUTES = ['buying', 'maint', 'doors', 'persons', 'lug_boot',
              'safety', 'class']
```

```
In [4]: data = pd.read_csv("data/car.data", names=ATTRIBUTES)
```

2. Discoverng and visualizing the data to gain insights

```
In [5]: data.head(10)
```

```
Out[5]:
```

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc
5	vhigh	vhigh	2	2	med	high	unacc
6	vhigh	vhigh	2	2	big	low	unacc
7	vhigh	vhigh	2	2	big	med	unacc
8	vhigh	vhigh	2	2	big	high	unacc
9	vhigh	vhigh	2	4	small	low	unacc

The attributes are Ordinal (Categorical)

```
In [6]: for i in data.columns:
        print("{} attribute has {} values".format(i, data[i].unique()))
```

```
buying attribute has ['vhigh' 'high' 'med' 'low'] values
maint attribute has ['vhigh' 'high' 'med' 'low'] values
doors attribute has ['2' '3' '4' '5more'] values
persons attribute has ['2' '4' 'more'] values
lug_boot attribute has ['small' 'med' 'big'] values
safety attribute has ['low' 'med' 'high'] values
class attribute has ['unacc' 'acc' 'vgood' 'good'] values
```


The strata (subgroups) are well-distributed (that's rare) except for the class attribute.

```
In [7]: fig, axes = plt.subplots(nrows=2, ncols=3)

data['buying'].value_counts().plot(kind='bar', ax=axes[0,1])
data['maint'].value_counts().plot(kind='bar', ax=axes[1,1])
data['doors'].value_counts().plot(kind='bar', ax=axes[0,0])
data['persons'].value_counts().plot(kind='bar', ax=axes[1,0])
data['lug_boot'].value_counts().plot(kind='bar', ax=axes[0,2])
data['safety'].value_counts().plot(kind='bar', ax=axes[1,2])
```


```
fig.tight_layout() # Or equivalently, "plt.tight_layout()"

plt.show()
```



In [8]:

```
data['class'].hist()
plt.show()
```



In [9]:

```
data.describe()
```

Out[9]:

	buying	maint	doors	persons	lug_boot	safety	class
count	1728	1728	1728	1728	1728	1728	1728
unique	4	4	4	3	3	3	4
top	vhigh	vhigh	2	2	small	low	unacc
freq	432	432	432	576	576	576	1210

Nice! There is no null value in the dataset.

In [10]:


```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   buying      1728 non-null   object
1   maint       1728 non-null   object
2   doors       1728 non-null   object
3   persons     1728 non-null   object
4   lug_boot    1728 non-null   object
5   safety      1728 non-null   object
6   class       1728 non-null   object
dtypes: object(7)
memory usage: 94.6+ KB
```

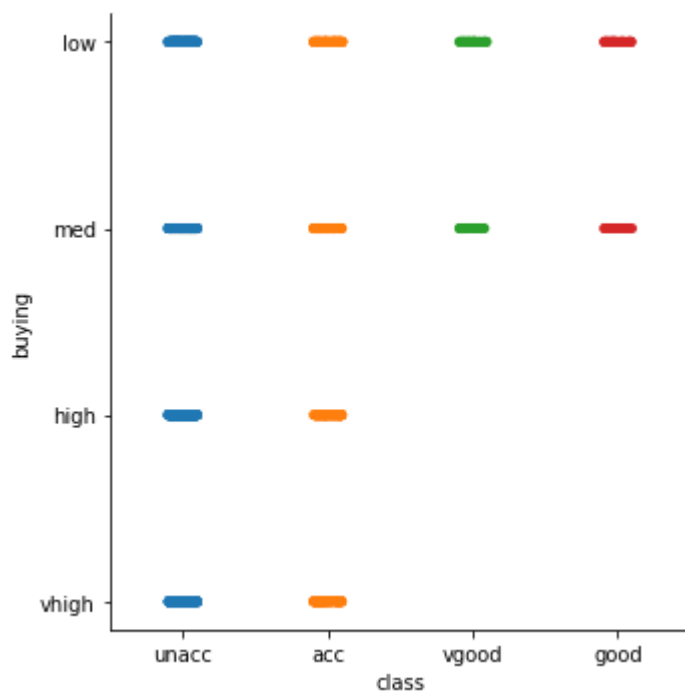
```
In [11]: data['buying'].hist()
```

```
Out[11]: <AxesSubplot:>
```



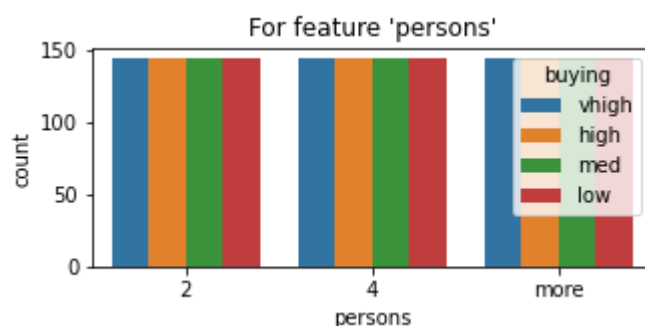
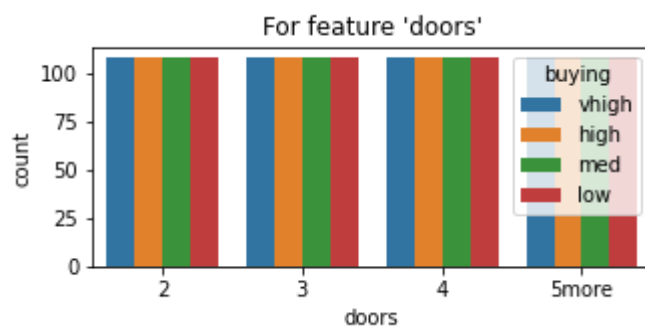
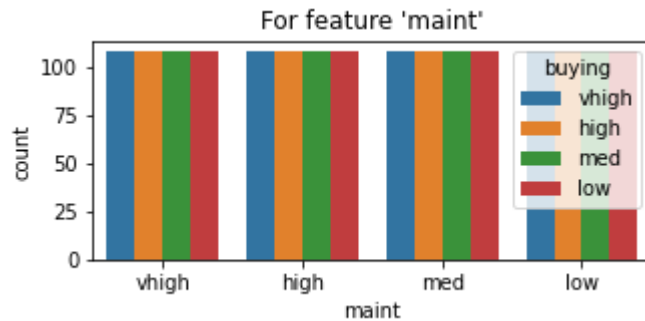
Negative correlation between class and buying

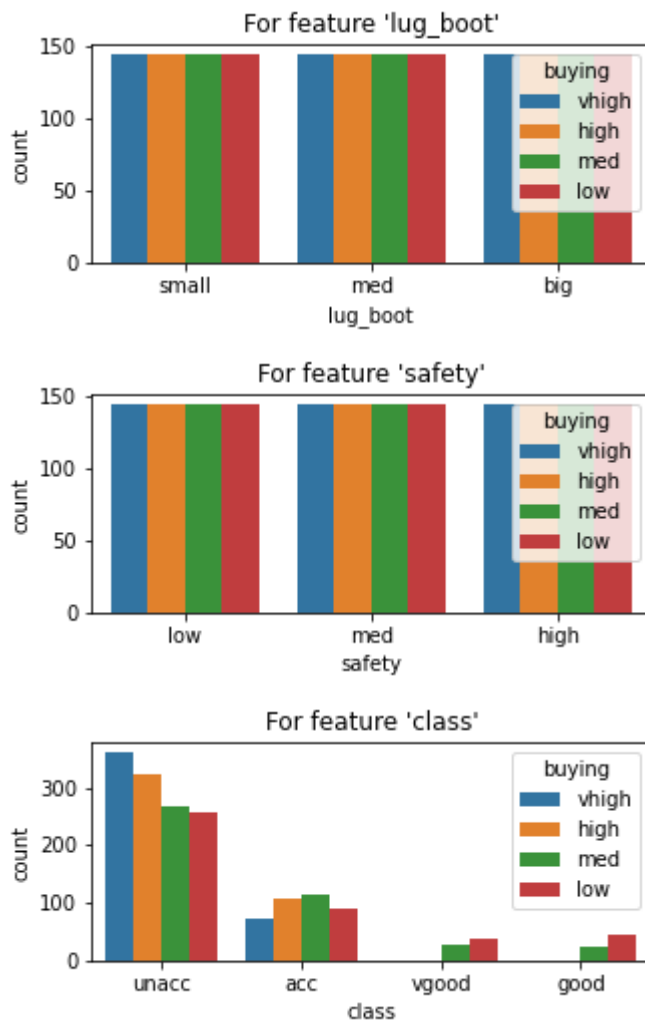
```
In [12]: sns.catplot(x="class", y="buying", data=data)
plt.show()
```



Visualizing the distribution of `buying` across attributes

```
In [13]: for i in data.columns[1:]:
plt.figure(figsize=(5,2))
plt.title("For feature '%s'"%i)
sns.countplot(data[i],hue=data['buying'])
```





vhigh and high buying price are associated with only unacc and acc classes.

From the EDA on the data we observe that there is no need to handle missing values, feature scaling, outliers handling and so on.

3. Preparing the data for Machine Learning Algorithms

Performing ordinal encoding since they give better measure of the distance between different categories of an ordinal attribute

```
In [14]: # ordinal mapping of the attributes
buying_mapping = {"vhigh": 3, "high": 2, "med": 1, 'low': 0}
maint_mapping = {"vhigh": 3, "high": 2, "med": 1, 'low': 0}
lug_boot_mapping = {"small": 0, "med": 1, "big": 2}
safety_mapping = {"low": 0, "med": 1, "high": 2}
class_mapping = {"unacc": 0, "acc": 1, "good": 2, "vgood": 3}
doors_mapping = {"2": 0, "3": 1, "4": 2, "5more": 3}
persons_mapping = {"2": 0, "3": 1, "more": 2}
```

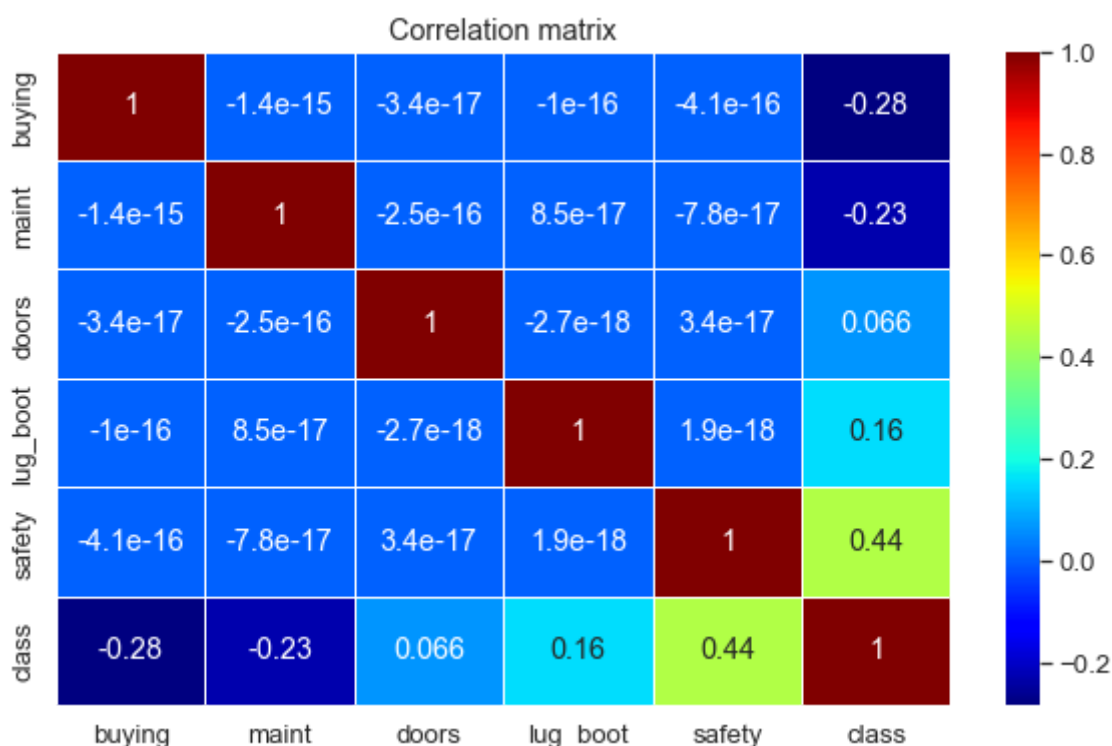
```
In [15]: # mapping the categories to the respective ordinal value for each attribute
trans_data = data.replace({"buying": buying_mapping})
trans_data = trans_data.replace({"maint": maint_mapping})
```

```
trans_data = trans_data.replace({"lug_boot": lug_boot_mapping})
trans_data = trans_data.replace({"safety": safety_mapping})
trans_data = trans_data.replace({"class": class_mapping})
trans_data = trans_data.replace({"persons": persons_mapping})
trans_data = trans_data.replace({"doors": doors_mapping})
```

Now we have the data in numeric format let's create the correlation matrix.

In [16]:

```
plt.figure(figsize=(10,6))
sns.set(font_scale=1.2)
sns.heatmap(trans_data.corr(),annot=True,
cmap='jet',linewidth=0.5)
plt.title('Correlation matrix');
```



class and maint are highly correlated (negatively)

buying and class are highly correlated (negatively)

Testing the initial assumption

Let's look at the correlation matrix again:

In [17]:

```
corr_matrix = trans_data.corr()
corr_matrix['buying'].sort_values(ascending=False)
```

Out[17]:

```
buying      1.000000e+00
doors       -3.433468e-17
lug_boot    -1.045866e-16
safety      -4.082114e-16
maint       -1.356939e-15
```

```
class      -2.827504e-01
Name: buying, dtype: float64
```


Attributes are negatively correlated with the `buying` attribute

This makes the initial assumptions regarding positive correlation between buying and {doors, lug_boot, safety, persons} FALSE.

Let's look into the distribution of *buying* categories on other attributes

In [18]:

```
sns.pairplot(trans_data, hue="buying", diag_kind="hist",
             height=1.5, palette="dark")
plt.show()
```



Splitting the data into training and testing

In [22]:

```
from sklearn.model_selection import train_test_split
```

In [23]:

```
X = trans_data.drop(['buying'], axis = 1)
y = trans_data['buying']
```

Following general convention - 80% training and 20% testing data, with random shuffling to prevent any potential selection bias


```
In [24]: # train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, shuffle=True, random_state=42)

# train_set['class'].value_counts()/len(train_set)
# test_set['class'].value_counts()/len(train_set)
```

4. Training, evaluating and error analysis on the training set

```
In [25]: # models
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeRegressor
from sklearn.neural_network import MLPClassifier
from sklearn import tree

# ensemble models
from sklearn.ensemble import RandomForestClassifier

# metrics
# from sklearn.metrics import log_loss
from sklearn.model_selection import cross_val_score,
learning_curve, GridSearchCV
from sklearn.metrics import confusion_matrix, mean_squared_error, \
mean_absolute_error, classification_report, precision_score,
recall_score, f1_score, \
matthews_corrcoef, cohen_kappa_score, roc_auc_score
```

```
In [26]: # good to test on various metrics for the multi-classification
task
def cal_metrics(y_test, y_pred):
    print(classification_report(y_test, y_pred))
    print('Recall: %.3f' % recall_score(y_test, y_pred,
average='micro'))
    print('F1: %.3f' % f1_score(y_test, y_pred, average='micro'))
    print("matthews_corrcoef: %.3f" % matthews_corrcoef(y_test,
y_pred))
    print("cohen_kappa_score: %.3f" % cohen_kappa_score(y_test,
y_pred))
```

```

    print("Mean squared error (MSE): %.3f" %
mean_squared_error(y_test, y_pred))
    print("Mean absolute error (MAE): %.3f" %
mean_absolute_error(y_test, y_pred))

# confusion matrix for error analysis
def create_confusion_matrix(y_test, y_pred, model_type):
    # Plot confusion matrix.
    cm_matrix = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(4,4))
    sns.set(font_scale=1.4)
    sns.heatmap(cm_matrix, annot=True, cbar=True,
cmap='jet',linewidth=0.5,fmt="d")
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.title('Confusion Matrix for {}'.format(model_type));

```

4.1. Logistic Regression

In [27]:

```

# Initialize a Logistic Regression classifier.
logreg = LogisticRegression(solver='saga', multi_class='auto',
random_state=42, n_jobs=-1)
logreg.fit(X_train, y_train)

# cross validation evaluation
logreg_cv = cross_val_score(logreg, X_train, y_train, cv=20)

```

In [28]:

```

# CV Accuracy
print('CV Accuracy: %.3f' % logreg_cv.mean())

```

CV Accuracy: 0.328

Predictions and error analysis

In [29]:

```

logreg_pred = logreg.predict(X_test)
cal_metrics(y_test, logreg_pred)
print('\n')
create_confusion_matrix(y_test, logreg_pred, "Logistic
Regression")

```

	precision	recall	f1-score	support
0	0.36	0.43	0.39	83
1	0.21	0.19	0.20	77
2	0.26	0.13	0.17	92

	predict_car_price			
	3	0.37	0.50	0.42
				94
accuracy			0.32	346
macro avg	0.30	0.31	0.30	346
weighted avg	0.30	0.32	0.30	346

Recall: 0.318

F1: 0.318


matthews_corrcoef: 0.091

cohen_kappa_score: 0.089

Mean squared error (MSE): 1.997

Mean absolute error (MAE): 1.049

Confusion Matrix for Logistic Regression



4.2. Logistic Regression - Grid Search

In [192...

```
# Hyperparameters to be checked.
parameters = {'C':[0.0001,0.001, 0.01, 1, 0.1, 10, 100, 1000],
              'penalty':['none','l2'],
              'solver':['lbfgs','sag','saga','newton-cg']}

# Logistic Regression classifier.
default_logreg=LogisticRegression(multi_class='auto',
random_state=42, n_jobs=-1)

# GridSearchCV estimator.
gs_logreg = GridSearchCV(default_logreg, parameters, cv=10,
verbose=1)

# Train the GridSearchCV estimator and search for the best
parameters.
gs_logreg.fit(X_train,y_train)
```

Fitting 10 folds for each of 64 candidates, totalling 640 fits

```
Out[192]: GridSearchCV(cv=10, estimator=LogisticRegression(n_jobs=-1, random_state=42),
               param_grid={'C': [0.0001, 0.001, 0.01, 1, 0.1, 10, 100, 1000],
                           'penalty': ['none', 'l2'],
                           'solver': ['lbfgs', 'sag', 'saga', 'newton-cg']},
               verbose=1)
```

```
In [200]: # Best parameters.
print("Best Logistic Regression Parameters:
      {}".format(gs_logreg.best_params_))

# Cross validation accuracy for the best parameters.
print('CV accuracy: %0.3f' % gs_logreg.best_score_)

# Accuracy: 1 is perfect prediction.
print('Accuracy: %0.3f' % (gs_logreg.score(X_test, y_test)))
```

Best Logistic Regression Parameters: {'C': 0.1, 'penalty': 'l2', 'solver': 'lbfgs'}

CV accuracy: 0.329

Accuracy: 0.327

4.3. K-Nearest Neighbours

```
In [60]: # Initialize a decision tree estimator.
knn = KNeighborsClassifier(n_jobs=-1, n_neighbors=300)
# Train the estimator.
knn.fit(X_train, y_train)

# cross validation evaluation
knn_cv = cross_val_score(knn, X_train, y_train, cv=20)
```

```
In [62]: # CV Accuracy
print('CV Accuracy: %0.3f' % knn_cv.mean())
```

CV Accuracy: 0.313

Predictions and error analysis

```
In [64]: knn_pred = knn.predict(X_test)
cal_metrics(y_test, knn_pred)
print('\n')
create_confusion_matrix(y_test, knn_pred, "K-Nearest Neighbours")
```

	precision	recall	f1-score	support
0	0.39	0.24	0.30	83
1	0.18	0.25	0.21	77

				predict_car_price	
	2	0.32	0.32	0.32	92
	3	0.35	0.37	0.36	94
	accuracy			0.30	346
	macro avg	0.31	0.29	0.30	346
	weighted avg	0.31	0.30	0.30	346

Recall: 0.298

F1: 0.298

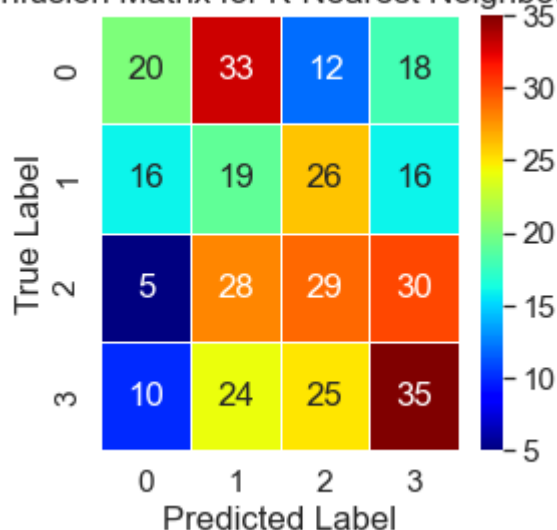
matthews_corrcoef: 0.063

cohen_kappa_score: 0.063

Mean squared error (MSE): 1.844

Mean absolute error (MAE): 1.029

Confusion Matrix for K-Nearest Neighbours



4.4. Decision Trees

```
In [85]: # Initialize a decision tree estimator.
tr = tree.DecisionTreeClassifier(max_depth=4, criterion='entropy',
random_state=42)
# Train the estimator.
tr.fit(X_train, y_train)

# cross validation evaluation
tr_cv = cross_val_score(tr, X_train, y_train, cv=20)
```

```
In [86]: # CV Accuracy
print('CV Accuracy: %.3f' % tr_cv.mean())
```

CV Accuracy: 0.311

Predictions and error analysis

```
In [88]: tr_pred = tr.predict(X_test)
cal_metrics(y_test, tr_pred)
```

```
print('\n')
create_confusion_matrix(y_test, tr_pred, "Decision trees")
```

	precision	recall	f1-score	support
0	0.25	0.58	0.35	83
1	0.22	0.26	0.24	77
2	0.20	0.04	0.07	92
3	0.57	0.26	0.35	94
accuracy			0.28	346
macro avg	0.31	0.28	0.25	346
weighted avg	0.32	0.28	0.25	346

Recall: 0.277

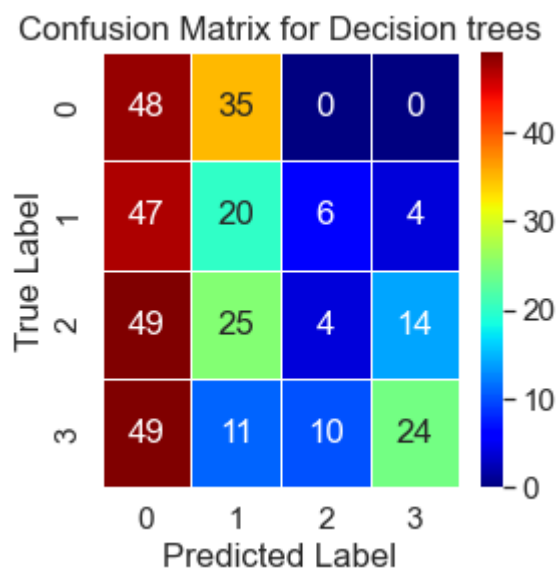
F1: 0.277

matthews_corrcoef: 0.055

cohen_kappa_score: 0.048

Mean squared error (MSE): 2.410

Mean absolute error (MAE): 1.191



4.5. Decision Trees - Grid Search

In [195...

```
# Hyperparameters to be checked.
parameters = {'criterion':['gini','entropy'],
              'max_depth':
[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20]
              }

# Default Decision tree estimator.
default_tr = tree.DecisionTreeClassifier(random_state=42)

# GridSearchCV estimator.
gs_tree = GridSearchCV(default_tr, parameters, cv=10,
```

```
n_jobs=-1,verbose=1)

# Train the GridSearchCV estimator and search for the best
parameters.

gs_tree.fit(X_train,y_train)
```

Fitting 10 folds for each of 40 candidates, totalling 400 fits

```
Out[195]: GridSearchCV(cv=10, estimator=DecisionTreeClassifier(random_state=42),
                    n_jobs=-1,
                    param_grid={'criterion': ['gini', 'entropy'],
                                'max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
                                                13, 14, 15, 16, 17, 18, 19, 20]},
                    verbose=1)
```

```
In [199]: # Best parameters.
print("Best Decision Tree Parameters:
{}".format(gs_tree.best_params_))

# Cross validation accuracy for the best parameters.
print('CV accuracy: %0.3f' % gs_tree.best_score_)

# Accuracy: 1 is perfect prediction.
print('Accuracy: %0.3f' % (gs_tree.score(X_test, y_test)))
```

Best Decision Tree Parameters: {'criterion': 'entropy', 'max_depth': 3}
 CV accuracy: 0.321
 Accuracy: 0.298

4.6. Multi Layered Perceptron

```
In [93]: mlp = MLPClassifier(hidden_layer_sizes=(10), max_iter=100,
                    random_state=42, shuffle=True, verbose=False)
mlp.fit(X_train, y_train)

# cross validation evaluation
mlp_cv = cross_val_score(mlp, X_train, y_train, cv=20)
```

```
In [94]: # CV Accuracy
print('CV Accuracy: %0.3f' % mlp_cv.mean())
```

CV Accuracy: 0.321

Predictions and error analysis

```
In [96]: mlp_pred = mlp.predict(X_test)
cal_metrics(y_test, mlp_pred)
print('\n')
```

```
create_confusion_matrix(y_test, mlp_pred, "Multi Layered
Perceptron")
```

	precision	recall	f1-score	support
0	0.50	0.41	0.45	83
1	0.23	0.30	0.26	77
2	0.25	0.16	0.20	92
3	0.36	0.45	0.40	94
accuracy			0.33	346
macro avg	0.33	0.33	0.33	346
weighted avg	0.33	0.33	0.33	346

Recall: 0.329

F1: 0.329


matthews_corrcoef: 0.107

cohen_kappa_score: 0.106

Mean squared error (MSE): 1.624

Mean absolute error (MAE): 0.954

Confusion Matrix for Multi Layered Perceptron



4.7. Multi Layered Perceptron - Grid Search

```
In [ ]: # Hyperparameters to be checked.
parameters = {'activation':['logistic','tanh','relu'],
              'solver': ['lbfgs','adam','sgd'],
              'alpha':10.0 ** -np.arange(1,3),
              'hidden_layer_sizes':[(5),(100),(3),(4),(3,1),
              (5,3)]}

# MLP estimator.
default_mlp = MLPClassifier(random_state=42)

# GridSearchCV estimator.
```



```
gs_mlp = GridSearchCV(default_mlp, parameters, cv=10,
n_jobs=-1,verbose=1)

# Train the GridSearchCV estimator and search for the best
parameters.

gs_mlp.fit(X_train,y_train)
```

In [198...

```
gs_mlp_pred = gs_mlp.predict(X_test)
cal_metrics(y_test, gs_mlp_pred)
print('\n')
create_confusion_matrix(y_test, gs_mlp_pred, "Multi Layered
Perceptron")
```

	precision	recall	f1-score	support
0	0.59	0.35	0.44	83
1	0.20	0.27	0.23	77
2	0.37	0.24	0.29	92
3	0.41	0.59	0.48	94
accuracy			0.37	346
macro avg	0.39	0.36	0.36	346
weighted avg	0.40	0.37	0.37	346

Recall: 0.367

F1: 0.367


matthews_corrcoef: 0.158

cohen_kappa_score: 0.154

Mean squared error (MSE): 1.569

Mean absolute error (MAE): 0.910

Confusion Matrix for Multi Layered Perceptron



In [190...

```
# Best parameters.
print("Best MLP Parameters: {}".format(gs_mlp.best_params_))
```

```
# Cross validation accuracy for the best parameters.
print('CV accuracy: %0.3f' % gs_mlp.best_score_)

# Accuracy: 1 is perfect prediction.
print('Accuracy: %0.3f' % (gs_mlp.score(X_test,y_test)))
```

Best MLP Parameters: {'activation': 'logistic', 'alpha': 0.01, 'hidden_layer_sizes': (3, 1), 'solver': 'lbfgs'}

CV accuracy: 0.353

Accuracy: 0.367

4.8. Ensembling - Random Forest

In [182...

```
# Initialize a Random Forest classifier.
rfc = RandomForestClassifier(n_jobs=-1,
                             min_samples_leaf=2,
                             max_depth=4,
                             n_estimators=10,
                             min_samples_split=5,
                             random_state=51)

rfc.fit(X_train, y_train)

# cross validation evaluation
rfc_cv = cross_val_score(rfc, X_train, y_train, cv=20)
```

In [183...

```
# CV Accuracy
print('CV Accuracy: %0.3f' % rfc_cv.mean())
```

CV Accuracy: 0.319

Predictions and error analysis


In [184...

```
rfc_pred = rfc.predict(X_test)
cal_metrics(y_test, rfc_pred)
print('\n')
create_confusion_matrix(y_test, rfc_pred, "Random Forest")
```

	precision	recall	f1-score	support
0	0.45	0.27	0.33	83
1	0.26	0.22	0.24	77
2	0.29	0.23	0.25	92
3	0.30	0.51	0.38	94
accuracy			0.31	346
macro avg	0.32	0.31	0.30	346
weighted avg	0.32	0.31	0.30	346

Recall: 0.312
 F1: 0.312
 matthews_corrcoef: 0.077
 cohen_kappa_score: 0.075
 Mean squared error (MSE): 1.974
 Mean absolute error (MAE): 1.043

Confusion Matrix for Random Forest



5. Selecting the best model

Neural network (Multi Layered perceptron) with Grid Search performed the best among other machine learning models tested.

With parameters: {'activation': 'logistic', 'alpha': 0.01,
 'hidden_layer_sizes': (3, 1), 'solver': 'lbfgs'}
 CV accuracy: 0.353
 Accuracy: 0.367

6. Conclusions: Presenting solutions

With data analysis and modelling I have following conclusions:

1. The best model with 36.7% performs far worse than the random chance (50% accuracy). Such models are NOT suitable for a business use case. Rather a data-driven approach needs to be applied rather than focusing on a model-driven approach i.e. finding the best model.
2. Clearly, with 36.7% accuracy the model didn't converge to an optimal loss.
3. Error analysis with confusion matrix reflects random guessing across the buying price categories.
4. The data failed my assumptions of the correlations. While my assumptions might be flawed but I believe that help from an expert in the required field can help curate the

data.

5. The data already has fewer features and even if I dropped the least important feature or correlated ones, then also the accuracy was reduced by around 1–2% . So, dropping a feature is not an option to reduce variance in a model.
6. Collecting other features (attributes) and scaling data can be helpful.